

In this Python notebook, I will be analyzing the dataset to see if there are any features or patterns I can leverage to better build my model, which aims to predict median house values. I will be using `numpy`, `pandas`, and `seaborn` to help me go through the data. For model implementation, I will be using XGBoost. The dataset used is the California housing data set, which can be found on Kaggle here: <https://www.kaggle.com/datasets/camnugent/california-housing-prices>

```
# Necessary imports
import numpy as np
import pandas as pd
import seaborn as sns
```

As this is a CSV file, I will first load the file into a `pandas` dataframe and will take a look at the format before deciding what to do next.

```
# Read CSV file into dataframe
df = pd.read_csv("housing.csv")
```

df

	longitude	latitude	housing_median_age	total_rooms
total_bedrooms \				
0	-122.23	37.88	41.0	880.0
129.0				
1	-122.22	37.86	21.0	7099.0
1106.0				
2	-122.24	37.85	52.0	1467.0
190.0				
3	-122.25	37.85	52.0	1274.0
235.0				
4	-122.25	37.85	52.0	1627.0
280.0				
...	...	...	...	...
...				
20635	-121.09	39.48	25.0	1665.0
374.0				
20636	-121.21	39.49	18.0	697.0
150.0				
20637	-121.22	39.43	17.0	2254.0
485.0				
20638	-121.32	39.43	18.0	1860.0
409.0				
20639	-121.24	39.37	16.0	2785.0
616.0				

	population	households	median_income	median_house_value \
0	322.0	126.0	8.3252	452600.0
1	2401.0	1138.0	8.3014	358500.0
2	496.0	177.0	7.2574	352100.0

3	558.0	219.0	5.6431	341300.0
4	565.0	259.0	3.8462	342200.0
...	...	...	...	...
20635	845.0	330.0	1.5603	78100.0
20636	356.0	114.0	2.5568	77100.0
20637	1007.0	433.0	1.7000	92300.0
20638	741.0	349.0	1.8672	84700.0
20639	1387.0	530.0	2.3886	89400.0

	ocean_proximity
0	NEAR BAY
1	NEAR BAY
2	NEAR BAY
3	NEAR BAY
4	NEAR BAY
...	...
20635	INLAND
20636	INLAND
20637	INLAND
20638	INLAND
20639	INLAND

[20640 rows x 10 columns]

The next step is to clean the data. There is potential for missing data, so the dataframe must be checked for any rows that are missing any data. If there are very few in comparison to the total number of rows (20640), then the rows will be entirely removed.

*# Figure out how many rows of missing data there is*

```
missing_rows = df[df.isnull().any(axis=1)]
```

missing\_rows

	longitude	latitude	housing_median_age	total_rooms
total_bedrooms \				
290	-122.16	37.77	47.0	1256.0
NaN				
341	-122.17	37.75	38.0	992.0
NaN				
538	-122.28	37.78	29.0	5154.0
NaN				
563	-122.24	37.75	45.0	891.0
NaN				
696	-122.10	37.69	41.0	746.0
NaN				
...	...	...	...	...
...				
20267	-119.19	34.20	18.0	3620.0
NaN				

```

20268    -119.18    34.19                19.0    2393.0
NaN
20372    -118.88    34.17                15.0    4260.0
NaN
20460    -118.75    34.29                17.0    5512.0
NaN
20484    -118.72    34.28                17.0    3051.0
NaN

      population  households  median_income  median_house_value  \
290           570.0       218.0         4.3750         161900.0
341           732.0       259.0         1.6196          85100.0
538          3741.0      1273.0         2.5762        173400.0
563           384.0       146.0         4.9489        247100.0
696           387.0       161.0         3.9063        178400.0
...          ...         ...         ...         ...
20267        3171.0       779.0         3.3409        220500.0
20268        1938.0       762.0         1.6953        167400.0
20372        1701.0       669.0         5.1033        410700.0
20460        2734.0       814.0         6.6073        258100.0
20484        1705.0       495.0         5.7376        218600.0

      ocean_proximity
290          NEAR BAY
341          NEAR BAY
538          NEAR BAY
563          NEAR BAY
696          NEAR BAY
...          ...
20267        NEAR OCEAN
20268        NEAR OCEAN
20372        <1H OCEAN
20460        <1H OCEAN
20484        <1H OCEAN

[207 rows x 10 columns]

```

There are only 207 rows with missing data, so it should be okay to remove them.

```

df = df.dropna()

df
      longitude  latitude  housing_median_age  total_rooms
total_bedrooms  \
0          -122.23    37.88                41.0         880.0
129.0
1          -122.22    37.86                21.0        7099.0
1106.0
2          -122.24    37.85                52.0        1467.0

```

```

190.0
3      -122.25      37.85      52.0      1274.0
235.0
4      -122.25      37.85      52.0      1627.0
280.0
...      ...      ...      ...      ...
...
20635      -121.09      39.48      25.0      1665.0
374.0
20636      -121.21      39.49      18.0      697.0
150.0
20637      -121.22      39.43      17.0      2254.0
485.0
20638      -121.32      39.43      18.0      1860.0
409.0
20639      -121.24      39.37      16.0      2785.0
616.0

      population      households      median_income      median_house_value \
0          322.0          126.0          8.3252          452600.0
1          2401.0          1138.0          8.3014          358500.0
2           496.0           177.0          7.2574          352100.0
3           558.0           219.0          5.6431          341300.0
4           565.0           259.0          3.8462          342200.0
...      ...      ...      ...      ...
20635          845.0          330.0          1.5603          78100.0
20636          356.0          114.0          2.5568          77100.0
20637         1007.0          433.0          1.7000          92300.0
20638          741.0          349.0          1.8672          84700.0
20639         1387.0          530.0          2.3886          89400.0

      ocean_proximity
0          NEAR BAY
1          NEAR BAY
2          NEAR BAY
3          NEAR BAY
4          NEAR BAY
...      ...
20635          INLAND
20636          INLAND
20637          INLAND
20638          INLAND
20639          INLAND

[20433 rows x 10 columns]

```

This will one-hot encode the ocean proximity column so that it can be better analyzed by the model (generally numeric values are preferred).

```
df = pd.get_dummies(df, columns=["ocean_proximity"], drop_first=True)
```

```
df
```

	longitude	latitude	housing_median_age	total_rooms
total_bedrooms \				
0	-122.23	37.88	41.0	880.0
129.0				
1	-122.22	37.86	21.0	7099.0
1106.0				
2	-122.24	37.85	52.0	1467.0
190.0				
3	-122.25	37.85	52.0	1274.0
235.0				
4	-122.25	37.85	52.0	1627.0
280.0				
...	...	...	...	...
...				
20635	-121.09	39.48	25.0	1665.0
374.0				
20636	-121.21	39.49	18.0	697.0
150.0				
20637	-121.22	39.43	17.0	2254.0
485.0				
20638	-121.32	39.43	18.0	1860.0
409.0				
20639	-121.24	39.37	16.0	2785.0
616.0				

	population	households	median_income	median_house_value \
0	322.0	126.0	8.3252	452600.0
1	2401.0	1138.0	8.3014	358500.0
2	496.0	177.0	7.2574	352100.0
3	558.0	219.0	5.6431	341300.0
4	565.0	259.0	3.8462	342200.0
...	...	...	...	...
20635	845.0	330.0	1.5603	78100.0
20636	356.0	114.0	2.5568	77100.0
20637	1007.0	433.0	1.7000	92300.0
20638	741.0	349.0	1.8672	84700.0
20639	1387.0	530.0	2.3886	89400.0

	ocean_proximity_INLAND	ocean_proximity_ISLAND \
0	False	False
1	False	False
2	False	False
3	False	False
4	False	False
...	...	...
20635	True	False

20636	True	False
20637	True	False
20638	True	False
20639	True	False

	ocean_proximity_NEAR BAY	ocean_proximity_NEAR OCEAN
0	True	False
1	True	False
2	True	False
3	True	False
4	True	False
...	...	...
20635	False	False
20636	False	False
20637	False	False
20638	False	False
20639	False	False

[20433 rows x 13 columns]

This is the initial model that will be used as a baseline:

```

from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score
from sklearn.model_selection import train_test_split
from xgboost import XGBRegressor

def run_model(df):
    x = df.drop(columns=["median_house_value"])
    y = df["median_house_value"]

    # Split the data into training and testing sets
    X_train, X_test, y_train, y_test = train_test_split(x, y,
test_size=0.2, random_state=419)

    # Initialize the XGBoost regressor
    original_model = XGBRegressor(objective="reg:squarederror",
n_estimators=1000, learning_rate=0.01, max_depth=7)

    # Train the model
    original_model.fit(X_train, y_train)

    # Predict on the test set
    y_pred = original_model.predict(X_test)

    # Evaluate the model
    mse = mean_squared_error(y_test, y_pred)
    mae = mean_absolute_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)

```

```

print(f"Mean Squared Error: ${mse:.2f}")
print(f"Mean Absolute Error: ${mae:.2f}")
print(f"R^2 Score: {r2}")

# Medical tolerance (15%)
tolerance = 0.15

# Calculate accuracy
accuracy = ((abs(y_pred - y_test) / y_test) <= tolerance).mean()
print(f"Accuracy within {tolerance * 100}% tolerance:
{accuracy:.2f}")

return {"model": original_model, "mae": mae, "mse": mse, "r^2":
r2, "accuracy": accuracy}

original_data = run_model(df)

Mean Squared Error: $2267413855.57
Mean Absolute Error: $31261.91
R^2 Score: 0.832266042287044
Accuracy within 15.0% tolerance: 59.70%

```

Data visualization is necessary for identifying what parameters are desirable for the model and if there are any data distributions to take care of (whether the data is very dense or spread apart and whether there are any significant outliers that can skew the data).

Since median house value is one of the columns in the data, comparing the other parameters against the median house value will be good for checking for correlation. First, we will check the importance of the different features.

```

columns = df.drop(columns=["median_house_value"]).columns
column_importance = dict(zip(columns, [float(val) for val in
original_data["model"].feature_importances_]))

sorted_col_importance = dict(sorted(column_importance.items(),
key=lambda item: item[1], reverse=True))

sorted_col_importance

{'ocean_proximity_INLAND': 0.5352905988693237,
'median_income': 0.22805730998516083,
'ocean_proximity_NEAR OCEAN': 0.038734279572963715,
'ocean_proximity_ISLAND': 0.03855923190712929,
'latitude': 0.029591072350740433,
'longitude': 0.027198996394872665,
'housing_median_age': 0.024348242208361626,
'population': 0.019737565889954567,
'total_bedrooms': 0.019181597977876663,
'ocean_proximity_NEAR BAY': 0.016545208171010017,

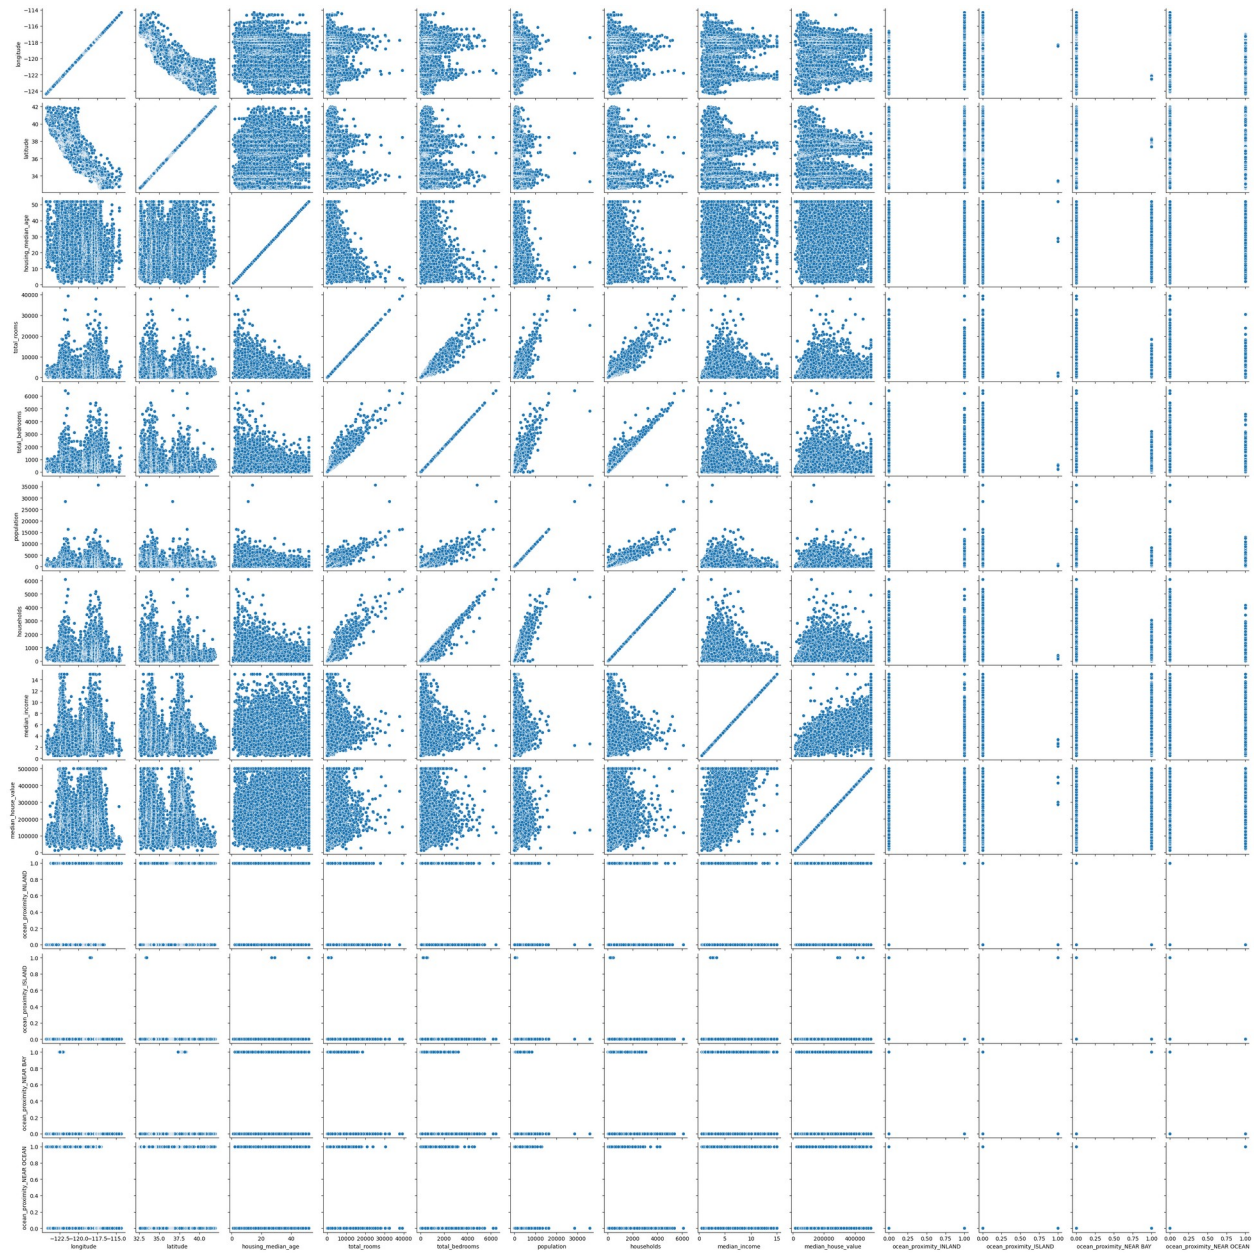
```

```
'households': 0.013357114978134632,  
'total_rooms': 0.009398694150149822}
```

We can see that `ocean_proximity_INLAND` and `median_income` are the two top features while `households` and `total_rooms` are not particularly helpful. Before removing or adding features, let's compare the categories with one another with a pairplot and a correlation heat map.


```
pairplot = sns.pairplot(df, diag_kind=None)  
pairplot  
<seaborn.axisgrid.PairGrid at 0x1bcfe42d940>
```





```
sns.heatmap(df.drop(columns="median_house_value").corr(), annot=True,
cmap="coolwarm", fmt=".2f")
```

<Axes: >



`total_rooms`, `total_bedrooms`, `population`, and `households` are all very correlated with one another ( $> 0.90$ ), and are all in the bottom five important features:

- `population`: 0.019737565889954567
- `total_bedrooms`: 0.019181597977876663
- `households`: 0.013357114978134632
- `total_rooms`: 0.009398694150149822

It may be wise to make new features that represent other attributes that may influence house value. Population density (`population / households`) could be useful along with rooms per house (average house size, `total_rooms / households`)

```
df_v2 = df.copy()

df_v2["pop_density"] = df_v2["population"] / df_v2["households"]
df_v2["house_density"] = df_v2["total_rooms"] / df_v2["households"]

df_v2
```

	longitude	latitude	housing_median_age	total_rooms
total_bedrooms \				
0	-122.23	37.88	41.0	880.0
129.0				
1	-122.22	37.86	21.0	7099.0
1106.0				
2	-122.24	37.85	52.0	1467.0
190.0				
3	-122.25	37.85	52.0	1274.0
235.0				
4	-122.25	37.85	52.0	1627.0
280.0				
...	...	...	...	...
...				
20635	-121.09	39.48	25.0	1665.0
374.0				
20636	-121.21	39.49	18.0	697.0
150.0				
20637	-121.22	39.43	17.0	2254.0
485.0				
20638	-121.32	39.43	18.0	1860.0
409.0				
20639	-121.24	39.37	16.0	2785.0
616.0				
	population	households	median_income	median_house_value \
0	322.0	126.0	8.3252	452600.0
1	2401.0	1138.0	8.3014	358500.0
2	496.0	177.0	7.2574	352100.0
3	558.0	219.0	5.6431	341300.0
4	565.0	259.0	3.8462	342200.0
...	...	...	...	...
20635	845.0	330.0	1.5603	78100.0
20636	356.0	114.0	2.5568	77100.0
20637	1007.0	433.0	1.7000	92300.0
20638	741.0	349.0	1.8672	84700.0
20639	1387.0	530.0	2.3886	89400.0
	ocean_proximity_INLAND	ocean_proximity_ISLAND \		
0	False	False		
1	False	False		
2	False	False		
3	False	False		
4	False	False		
...	...	...		
20635	True	False		
20636	True	False		
20637	True	False		
20638	True	False		
20639	True	False		

	ocean_proximity_NEAR	BAY	ocean_proximity_NEAR	OCEAN
pop_density \				
0		True		False
2.555556				
1		True		False
2.109842				
2		True		False
2.802260				
3		True		False
2.547945				
4		True		False
2.181467				
...		...		...
...				
20635		False		False
2.560606				
20636		False		False
3.122807				
20637		False		False
2.325635				
20638		False		False
2.123209				
20639		False		False
2.616981				

	house_density
0	6.984127
1	6.238137
2	8.288136
3	5.817352
4	6.281853
...	...
20635	5.045455
20636	6.114035
20637	5.205543
20638	5.329513
20639	5.254717

[20433 rows x 15 columns]

Now we try the model on the new data.

```
v2_data = run_model(df_v2)
```

Mean Squared Error: \$2189646288.23  
Mean Absolute Error: \$30181.74  
R^2 Score: 0.838018967285961  
Accuracy within 15.0% tolerance: 61.46%

From here, we can try to see if there is improvement on dropping any features whose importance is too few. We will check the feature importance again and decide on dropping features for the v3 model.

```
columns = df_v2.drop(columns=["median_house_value"]).columns
column_importance = dict(zip(columns, [float(val) for val in
v2_data["model"].feature_importances_]))

sorted_col_importance = dict(sorted(column_importance.items(),
key=lambda item: item[1], reverse=True))

sorted_col_importance
{'ocean_proximity_INLAND': 0.5162120461463928,
 'median_income': 0.25450021028518677,
 'pop_density': 0.05797304958105087,
 'ocean_proximity_ISLAND': 0.024760089814662933,
 'housing_median_age': 0.024340640753507614,
 'latitude': 0.02215234749019146,
 'longitude': 0.020298369228839874,
 'ocean_proximity_NEAR OCEAN': 0.018944622948765755,
 'ocean_proximity_NEAR BAY': 0.013564365915954113,
 'house_density': 0.012222538702189922,
 'total_bedrooms': 0.010043910704553127,
 'households': 0.009095135144889355,
 'population': 0.008325847797095776,
 'total_rooms': 0.007566750515252352}
```

As `pop_density` and `house_density` both out-weigh the four highly correlated, low importance features, for v3, we can drop `total_bedrooms`, `households`, `population`, and `total_rooms`.

```
df_v3 = df_v2.copy()

# Drop unimportant features
unimportant_features = ["total_bedrooms", "households", "population",
"total_rooms"]
df_v3 = df_v3.drop(columns=unimportant_features)
```

df\_v3

	longitude	latitude	housing_median_age	median_income	\
0	-122.23	37.88	41.0	8.3252	
1	-122.22	37.86	21.0	8.3014	
2	-122.24	37.85	52.0	7.2574	
3	-122.25	37.85	52.0	5.6431	
4	-122.25	37.85	52.0	3.8462	
...	...	...	...	...	
20635	-121.09	39.48	25.0	1.5603	
20636	-121.21	39.49	18.0	2.5568	

20637	-121.22	39.43	17.0	1.7000
20638	-121.32	39.43	18.0	1.8672
20639	-121.24	39.37	16.0	2.3886

	median_house_value	ocean_proximity_INLAND
ocean_proximity_ISLAND \		
0	452600.0	False
False		
1	358500.0	False
False		
2	352100.0	False
False		
3	341300.0	False
False		
4	342200.0	False
False		
...	...	...
...		
20635	78100.0	True
False		
20636	77100.0	True
False		
20637	92300.0	True
False		
20638	84700.0	True
False		
20639	89400.0	True
False		

	ocean_proximity_NEAR BAY	ocean_proximity_NEAR OCEAN
pop_density \		
0	True	False
2.555556		
1	True	False
2.109842		
2	True	False
2.802260		
3	True	False
2.547945		
4	True	False
2.181467		
...	...	...
..		
20635	False	False
2.560606		
20636	False	False
3.122807		
20637	False	False
2.325635		

20638	False	False
2.123209		
20639	False	False
2.616981		

	house_density
0	6.984127
1	6.238137
2	8.288136
3	5.817352
4	6.281853
...	...
20635	5.045455
20636	6.114035
20637	5.205543
20638	5.329513
20639	5.254717

[20433 rows x 11 columns]

Now let's try running the model to see what happens.

```
v3_data = run_model(df_v3)

Mean Squared Error: $2135919585.43
Mean Absolute Error: $29871.71
R^2 Score: 0.84199344793648
Accuracy within 15.0% tolerance: 61.83%
```




There was slight improvement but not much at all. Let's visualize the data now and see if there are any outliers to take care of.

```
features_to_compare =
df_v3.drop(columns=["median_house_value"]).columns

# Create the pairplot
sns.pairplot(df_v3, y_vars=features_to_compare,
x_vars="median_house_value", height=3, aspect=1, kind="scatter")

<seaborn.axisgrid.PairGrid at 0x1bd02880170>
```







From this, we can see that there are clear outliers:

- `ocean_proximity_ISLAND`: 4 outliers (only ones at 1.0). This is a one-hot encoded value though, so it is better to drop this feature altogether
- `pop_density`: 4 outliers (> 200). These points will just be removed.
- `house_density`: 10 outliers (> 40). These points will just be removed.

```
cleaned_v3 = df_v3.copy()
```

```
# Remove outliers
```

```
cleaned_v3 = cleaned_v3.drop(columns="ocean_proximity_ISLAND")
```

```
cleaned_v3 = cleaned_v3[cleaned_v3["pop_density"] <= 200]
```

```
cleaned_v3 = cleaned_v3[cleaned_v3["house_density"] <= 40]
```

```
cleaned_v3
```

	longitude	latitude	housing_median_age	median_income	\
0	-122.23	37.88	41.0	8.3252	
1	-122.22	37.86	21.0	8.3014	
2	-122.24	37.85	52.0	7.2574	
3	-122.25	37.85	52.0	5.6431	
4	-122.25	37.85	52.0	3.8462	
...	...	...	...	...	
20635	-121.09	39.48	25.0	1.5603	
20636	-121.21	39.49	18.0	2.5568	
20637	-121.22	39.43	17.0	1.7000	
20638	-121.32	39.43	18.0	1.8672	
20639	-121.24	39.37	16.0	2.3886	

	median_house_value	ocean_proximity_INLAND
ocean_proximity_NEAR BAY	\	
0	452600.0	False
True		
1	358500.0	False
True		
2	352100.0	False
True		
3	341300.0	False
True		
4	342200.0	False
True		
...	...	...
...		
20635	78100.0	True
False		
20636	77100.0	True
False		
20637	92300.0	True
False		
20638	84700.0	True

False			
20639	89400.0		True
False			

	ocean_proximity_NEAR OCEAN	pop_density	house_density
0	False	2.555556	6.984127
1	False	2.109842	6.238137
2	False	2.802260	8.288136
3	False	2.547945	5.817352
4	False	2.181467	6.281853
...	...	...	...
20635	False	2.560606	5.045455
20636	False	3.122807	6.114035
20637	False	2.325635	5.205543
20638	False	2.123209	5.329513
20639	False	2.616981	5.254717

[20418 rows x 10 columns]

Let's run the model once more to see if the cleaning helped any.

```
cleaned_v3_data = run_model(cleaned_v3)
```

Mean Squared Error: \$2036587664.55  
Mean Absolute Error: \$29048.12  
R<sup>2</sup> Score: 0.8492593892670817  
Accuracy within 15.0% tolerance: 62.81%

There was again a very slight improvement, but mostly insignificant. Let's check the feature importances now:

```
columns = cleaned_v3.drop(columns=["median_house_value"]).columns
column_importance = dict(zip(columns, [float(val) for val in
cleaned_v3_data["model"].feature_importances_]))

sorted_col_importance = dict(sorted(column_importance.items(),
key=lambda item: item[1], reverse=True))

sorted_col_importance
```

```
{'ocean_proximity_INLAND': 0.5880390405654907,
'median_income': 0.2376914918422699,
'pop_density': 0.053230006247758865,
'ocean_proximity_NEAR OCEAN': 0.02432853728532791,
'latitude': 0.023902567103505135,
'housing_median_age': 0.02278636209666729,
'longitude': 0.02268434502184391,
'house_density': 0.014687780290842056,
'ocean_proximity_NEAR BAY': 0.01264976616948843}
```

```

from sklearn.model_selection import GridSearchCV
from xgboost import XGBRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error,
r2_score

# Define features and target
x = cleaned_v3.drop(columns=["median_house_value"])
y = cleaned_v3["median_house_value"]

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(x, y,
test_size=0.2, random_state=419)

# Define the parameter grid
param_grid = {
    "n_estimators": [100, 500, 1000],
    "learning_rate": [0.01, 0.05, 0.1],
    "max_depth": [3, 5, 7],
    "subsample": [0.6, 0.8, 1.0],
    "colsample_bytree": [0.6, 0.8, 1.0]
}

# Initialize the XGBoost regressor
xgb_model = XGBRegressor(objective="reg:squarederror",
random_state=419)

# Set up GridSearchCV
grid_search = GridSearchCV(estimator=xgb_model, param_grid=param_grid,
scoring="neg_mean_squared_error", cv=5, verbose=1)

# Fit the GridSearchCV to find the best parameters
grid_search.fit(X_train, y_train)

# Get the best parameters
print("Best Parameters:", grid_search.best_params_)

# Train the model with the best parameters
best_model = grid_search.best_estimator_

# Predict on the test set
y_pred = best_model.predict(X_test)

# Evaluate the optimized model
mse = mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Optimized Mean Squared Error: ${mse:.2f}")
print(f"Optimized Mean Absolute Error: ${mae:.2f}")
print(f"Optimized R^2 Score: {r2}")

```

```

# Accuracy within tolerance
tolerance = 0.15
accuracy = ((abs(y_pred - y_test) / y_test) <= tolerance).mean()
print(f"Accuracy within {tolerance * 100}% tolerance: {accuracy:.2%}")

Fitting 5 folds for each of 243 candidates, totalling 1215 fits
Best Parameters: {'colsample_bytree': 0.6, 'learning_rate': 0.05,
'max_depth': 7, 'n_estimators': 1000, 'subsample': 1.0}
Optimized Mean Squared Error: $1843834097.42
Optimized Mean Absolute Error: $27400.77
Optimized R^2 Score: 0.8635262882255974
Accuracy within 15.0% tolerance: 66.26%

```

With this in mind, we now can adjust the `run_model()` function and focus more on data engineering since the improvements are modest. For graphing later, I will store the above results in a dictionary by running the new model on the same dataset, `cleaned_v3`.

```

# From above:
# Best Parameters: {'colsample_bytree': 0.6, 'learning_rate': 0.05,
'max_depth': 7, 'n_estimators': 1000, 'subsample': 1.0}

def run_better_model(df):
    x = df.drop(columns=["median_house_value"])
    y = df["median_house_value"]

    # Split the data into training and testing sets
    X_train, X_test, y_train, y_test = train_test_split(x, y,
test_size=0.2, random_state=419)

    # Initialize the XGBoost regressor
    original_model = XGBRegressor(objective="reg:squarederror",
n_estimators=1000, learning_rate=0.05, max_depth=7,
colsample_bytree=0.6, subsample=1.0)

    # Train the model
    original_model.fit(X_train, y_train)

    # Predict on the test set
    y_pred = original_model.predict(X_test)

    # Evaluate the model
    mse = mean_squared_error(y_test, y_pred)
    mae = mean_absolute_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)

    print(f"Mean Squared Error: ${mse:.2f}")
    print(f"Mean Absolute Error: ${mae:.2f}")
    print(f"R^2 Score: {r2}")

```

```

# Medical tolerance (15%)
tolerance = 0.15

# Calculate accuracy
accuracy = ((abs(y_pred - y_test) / y_test) <= tolerance).mean()
print(f"Accuracy within {tolerance * 100}% tolerance:
{accuracy:.2%}")

return {"model": original_model, "mae": mae, "mse": mse, "r^2":
r2, "accuracy": accuracy}

better_model_cleaned_v3_data = run_better_model(cleaned_v3)

Mean Squared Error: $1843834097.42
Mean Absolute Error: $27400.77
R^2 Score: 0.8635262882255974
Accuracy within 15.0% tolerance: 66.26%

```

There does seem to be a cap for the median house value (\$500,000), so let's see what happens if all values at that cap is removed (as it may skew the data).

```

df_v4 = cleaned_v3.copy()

# Remove the cap on house value
df_v4 = df_v4[df_v4["median_house_value"] < 500000]

df_v4

```

	longitude	latitude	housing_median_age	median_income	\
0	-122.23	37.88	41.0	8.3252	
1	-122.22	37.86	21.0	8.3014	
2	-122.24	37.85	52.0	7.2574	
3	-122.25	37.85	52.0	5.6431	
4	-122.25	37.85	52.0	3.8462	
...	...	...	...	...	
20635	-121.09	39.48	25.0	1.5603	
20636	-121.21	39.49	18.0	2.5568	
20637	-121.22	39.43	17.0	1.7000	
20638	-121.32	39.43	18.0	1.8672	
20639	-121.24	39.37	16.0	2.3886	

	median_house_value	ocean_proximity_INLAND
ocean_proximity_NEAR BAY \		
0	452600.0	False
True		
1	358500.0	False
True		
2	352100.0	False
True		
3	341300.0	False

```

True
4          342200.0          False
True
...          ...          ...
...
20635          78100.0          True
False
20636          77100.0          True
False
20637          92300.0          True
False
20638          84700.0          True
False
20639          89400.0          True
False

```

	ocean_proximity_NEAR OCEAN	pop_density	house_density
0	False	2.555556	6.984127
1	False	2.109842	6.238137
2	False	2.802260	8.288136
3	False	2.547945	5.817352
4	False	2.181467	6.281853
...	...	...	...
20635	False	2.560606	5.045455
20636	False	3.122807	6.114035
20637	False	2.325635	5.205543
20638	False	2.123209	5.329513
20639	False	2.616981	5.254717

```
[19434 rows x 10 columns]
```

```
v4_data = run_better_model(df_v4)
```

```
Mean Squared Error: $1652208484.28
```

```
Mean Absolute Error: $26686.92
```

```
R^2 Score: 0.8292834800864093
```

```
Accuracy within 15.0% tolerance: 65.35%
```

The accuracy went down a bit, but the mean absolute error is also a bit smaller. This is most likely due to there being a lot of cap-value houses that would have raised the MAE that are no longer in the dataset. It may be important to keep these values instead (and go with `cleaned_v3`).

Before graphing current results, let's try one more model, where we strictly keep the best features. Let's check the feature importance one more time and remove any significantly nonimportant features.

```

columns = cleaned_v3.drop(columns=["median_house_value"]).columns
column_importance = dict(zip(columns, [float(val) for val in
better_model_cleaned_v3_data["model"].feature_importances_]))

```

```
sorted_col_importance = dict(sorted(column_importance.items(),
key=lambda item: item[1], reverse=True))
```

```
sorted_col_importance
```

```
{'ocean_proximity_INLAND': 0.5031577348709106,
'median_income': 0.14498667418956757,
'latitude': 0.06451987475156784,
'ocean_proximity_NEAR OCEAN': 0.0592774972319603,
'pop_density': 0.05639267712831497,
'house_density': 0.05354450270533562,
'longitude': 0.05287262424826622,
'ocean_proximity_NEAR BAY': 0.03992283344268799,
'housing_median_age': 0.02532562054693699}
```

Using 0.05 as a cutoff, we can remove `ocean_proximity_NEAR BAY` and `housing_median_age`.

```
df_v5 = cleaned_v3.copy()
```

```
# Remove any feature with a score of < 0.05
```

```
df_v5 = df_v5.drop(columns=["ocean_proximity_NEAR BAY",
"housing_median_age"])
```

```
df_v5
```

	longitude	latitude	median_income	median_house_value \
0	-122.23	37.88	8.3252	452600.0
1	-122.22	37.86	8.3014	358500.0
2	-122.24	37.85	7.2574	352100.0
3	-122.25	37.85	5.6431	341300.0
4	-122.25	37.85	3.8462	342200.0
...	...	...	...	...
20635	-121.09	39.48	1.5603	78100.0
20636	-121.21	39.49	2.5568	77100.0
20637	-121.22	39.43	1.7000	92300.0
20638	-121.32	39.43	1.8672	84700.0
20639	-121.24	39.37	2.3886	89400.0

	ocean_proximity_INLAND	ocean_proximity_NEAR OCEAN	pop_density
0	False	False	2.555556
1	False	False	2.109842
2	False	False	2.802260
3	False	False	2.547945

4	False	False	2.181467
...	...	...	...
20635	True	False	2.560606
20636	True	False	3.122807
20637	True	False	2.325635
20638	True	False	2.123209
20639	True	False	2.616981

	house_density
0	6.984127
1	6.238137
2	8.288136
3	5.817352
4	6.281853
...	...
20635	5.045455
20636	6.114035
20637	5.205543
20638	5.329513
20639	5.254717

```
[20418 rows x 8 columns]
```

```
v5_data = run_better_model(df_v5)
```

Mean Squared Error: \$1981654355.69  
Mean Absolute Error: \$28515.40  
R^2 Score: 0.8533253475710927  
Accuracy within 15.0% tolerance: 64.57%

This is still worse than using `cleaned_v3` on the better model.

This is a tricky dataset as it is extremely large compared to what I project is to be the one used for the actual capstone project and it has real limitations (the target value has a value cap and is old data).

Overall, good improvements were made compared to the original dataset after cleaning, processing, and engineering. The graphs below showcase the models over time.

```
import matplotlib.pyplot as plt

all_data = [original_data, v2_data, v3_data, cleaned_v3_data,
better_model_cleaned_v3_data, v4_data, v5_data]
names = ["Original", "v2", "v3", "cleaned v3", "cleaned v3,\nnew
```



```

algo", "v4", "v5"]
colors = ["blue", "blue", "blue", "blue", "green", "blue", "blue"]

attributes = ["mae", "mse", "r^2", "accuracy"]
full_attributes = ["Mean Absolute Error", "Mean Squared Error", "r
squared", "Accuracy"]

# Make the bar chart
for attribute in attributes:
    # Create a new separate figure
    plt.figure()

    specific_data = []
    # Go through each set of data and collect the correct subset
    for data in all_data:
        specific_data.append(data[attribute])

    # Make the bar chart
    bars = plt.bar(names, specific_data, color=colors)

    # Rotate model names so they fit without overlap
    plt.xticks(rotation=90)


    # Set the y limits to better show the differences
    y_min = min(specific_data) * 0.95
    y_max = max(specific_data) * 1.05
    plt.ylim(y_min, y_max)


    # Labels
    plt.xlabel("Model Version")
    plt.ylabel(attribute)
    plt.title(f"Model Performance Measured by
{full_attributes[attributes.index(attribute)]}")

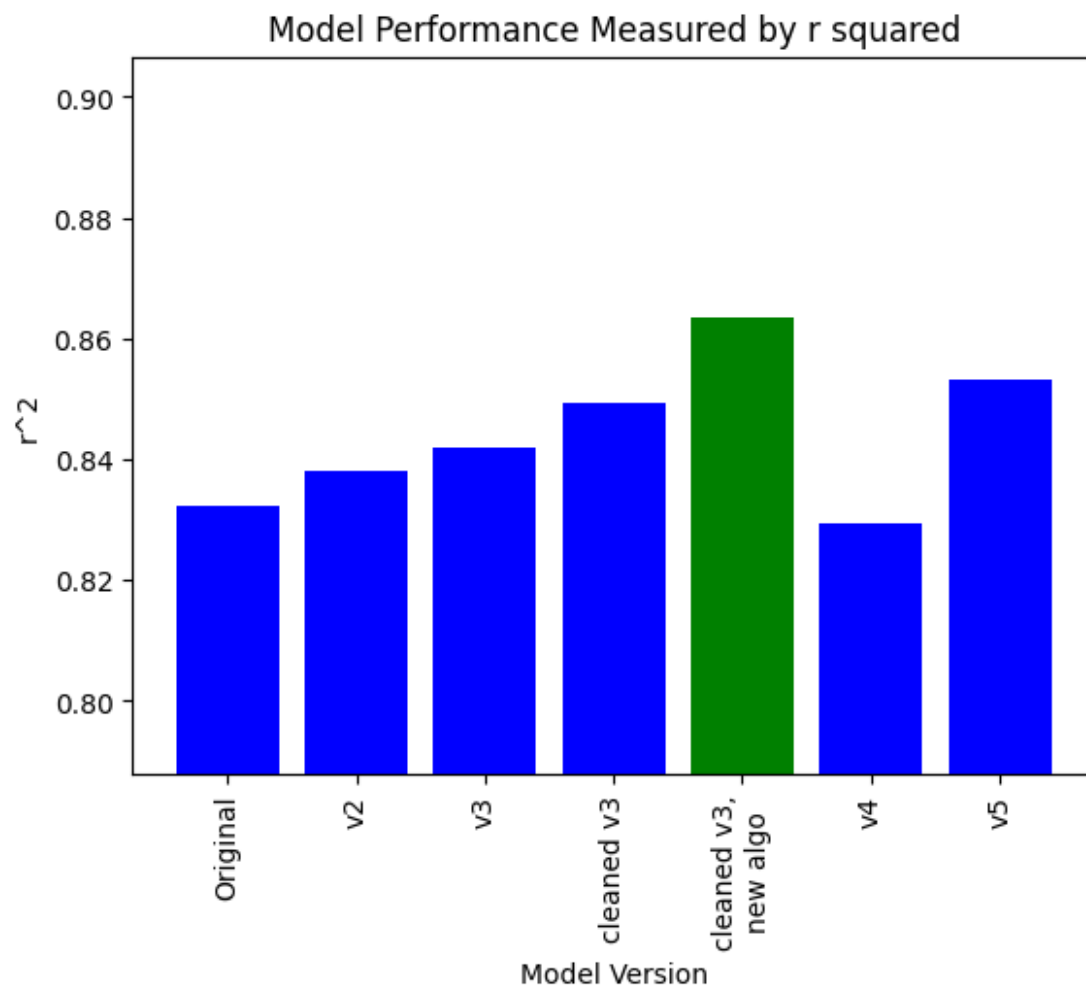
    plt.show()


```

Model Performance Measured by Mean Absolute Error









Below is a `pandas` table representation of the data.

```
data_breakdown = pd.DataFrame()

# Go through every index (every list has the same indices)
for index in range(len(names)):
    # Ensure name is properly formatted
    name = names[index]

    # For the sake of spacing, replace the newline with a regular
    space
    if "\n" in name:
        name = name.replace("\n", " ")

    # Format the values so they're more readable
    values = list(all_data[index].values())[1:]

    for value_index in range(len(values)):
        # Make accuracy a percentage
        if value_index == len(values)-1:
```

```

        values[value_index] *= 100

    values[value_index] = f"{values[value_index]:,.2f}"

    # Add percentage sign if on accuracy
    if value_index == len(values)-1:
        values[value_index] += "%"

    # Add a dollar sign if is a money error
    if value_index in (0, 1):
        values[value_index] = "$" + values[value_index]

    # Add the attributes
    data_breakdown[name] = values

# Set the labels for the table
data_breakdown.index = full_attributes

```

data\_breakdown

	Original	v2
v3 \		
Mean Absolute Error	\$31,261.91	\$30,181.74
\$29,871.71		
Mean Squared Error	\$2,267,413,855.57	\$2,189,646,288.23
\$2,135,919,585.43		
r squared	0.83	0.84
0.84		
Accuracy	59.70%	61.46%
61.83%		

	cleaned v3	cleaned v3, new algo \
Mean Absolute Error	\$29,048.12	\$27,400.77
Mean Squared Error	\$2,036,587,664.55	\$1,843,834,097.42
r squared	0.85	0.86
Accuracy	62.81%	66.26%

	v4	v5
Mean Absolute Error	\$26,686.92	\$28,515.40
Mean Squared Error	\$1,652,208,484.28	\$1,981,654,355.69
r squared	0.83	0.85
Accuracy	65.35%	64.57%